

A Scalable Two-Time-Scale Approach for Streaming Canonical Correlation and Generalized Eigenvalue Problems

Chris Junchi Li[◊]

Department of Electrical Engineering and Computer Sciences[◊]
University of California, Berkeley

October 8, 2024

Abstract

We propose TTS-CCA, a novel two-time-scale algorithm for streaming Canonical Correlation Analysis (CCA) and the generalized eigenvalue problem in high-dimensional settings. By extending Oja’s streaming PCA algorithm, we provide a provably efficient method for estimating principal generalized eigenvectors in the stochastic setting. The algorithm efficiently processes streaming data using a coupled iterative scheme, where one sequence evolves quickly while the other evolves slowly, allowing us to analyze convergence via the theory of two-time-scale stochastic approximation. Our main results establish a convergence rate of $O(1/n)$ under mild assumptions, significantly improving computational and memory efficiency compared to existing methods. In particular, our method only requires access to $O(d)$ memory, making it scalable to large-scale datasets. We also show that the proposed approach is robust to noise in the data, achieving optimal rates with respect to eigen-gap and dimensional dependence.

Keywords: Streaming CCA, Generalized Eigenvalue Problem, Two-Time-Scale Stochastic Approximation, Fast Convergence, High-Dimensional Data

1 Introduction

Canonical Correlation Analysis (CCA) and the Generalized Eigenvalue Problem (GEP) are two fundamental tools in machine learning and statistics, commonly used for feature extraction in applications such as regression [KF07], clustering [CKLS09], and classification [KM13]. Initially introduced by Hotelling [Hot36], CCA seeks to find a pair of linear transformations that maximize the correlation between two multidimensional random variables, making it a “correlation-aware” extension of Principal Component Analysis (PCA).

While CCA and GEP have proven invaluable in analyzing multi-view data, solving these problems efficiently in high-dimensional settings with streaming data remains a significant challenge. In particular, existing methods for CCA, although providing closed-form solutions, often suffer from high computational complexity, scaling as $O(d^3)$, which is prohibitive for large-scale datasets.

In this work, we address the challenges of streaming data and scalability by proposing TTS-CCA, a novel algorithm that extends Oja’s streaming PCA method to the stochastic CCA and GEP settings. Our approach leverages a two-time-scale stochastic approximation, where a fast-evolving sequence is used to update the estimate of the generalized eigenvectors, while a slower sequence ensures global convergence. We demonstrate that TTS-CCA achieves a convergence rate of $O(1/n)$ and requires only $O(d)$ memory, making it highly efficient for large-scale problems.

Our primary contributions include the introduction of TTS-CCA, a simple and globally convergent two-time-scale algorithm designed for solving the generalized eigenvalue problem and streaming

Canonical Correlation Analysis (CCA). We demonstrate that TTS-CCA achieves an optimal convergence rate of $O(1/n)$ under standard assumptions, matching the theoretical limits of stochastic approximation methods in non-convex optimization settings. Additionally, our method is computationally efficient and memory optimal, requiring only $O(d)$ storage, making it well-suited for large-scale and high-dimensional datasets.

Mathematically, Given access to samples $\{(x_i, y_i)_{i=1}^n\}$ of zero mean random variables $X, Y \in \mathbb{R}^d$ with an unknown joint distribution P_{XY} , CCA can be used to discover features expressing similarity or dissimilarity between X and Y . Formally, CCA aims to find a pair of vectors $u, v \in \mathbb{R}^d$ such that projections of X onto v and Y onto u are maximally correlated. In the population setting, the corresponding objective is given by:

$$\max v^\top \mathbb{E}[XY^\top]u \quad \text{s.t.} \quad v^\top \mathbb{E}[XX^\top]v = 1 \text{ and } u^\top \mathbb{E}[YY^\top]u = 1 \quad (1)$$

In the context of covariance matrices, the objective of the generalized eigenvalue problem is to obtain the direction u or $v \in \mathbb{R}^d$ maximizing discrepancy between X and Y and can be formulated as,

$$\arg \max_{v \neq 0} \frac{v^\top \mathbb{E}[XX^\top]v}{v^\top \mathbb{E}[YY^\top]v} \text{ and } \arg \max_{u \neq 0} \frac{u^\top \mathbb{E}[YY^\top]u}{u^\top \mathbb{E}[XX^\top]u} \quad (2)$$

More generally, given symmetric matrices A, B , with B positive definite, the objective of the principal generalized eigenvector problem is to obtain a unit norm vector w such that $Aw = \lambda Bw$ for λ maximal.

CCA and the generalized eigenvalue problem are intimately related. In fact, the CCA problem can be cast as a special case of the generalized eigenvalue problem by solving for u and v in the following objective:

$$\underbrace{\begin{pmatrix} 0 & \mathbb{E}[XY^\top] \\ \mathbb{E}[YX^\top] & 0 \end{pmatrix}}_A \begin{pmatrix} v \\ u \end{pmatrix} = \lambda \underbrace{\begin{pmatrix} \mathbb{E}[XX^\top] & 0 \\ 0 & \mathbb{E}[YY^\top] \end{pmatrix}}_B \begin{pmatrix} v \\ u \end{pmatrix} \quad (3)$$

The optimization problems underlying both CCA and the generalized eigenvector problem are non-convex in general. While they admit closed-form solutions, even in the offline setting a direct computation requires $\mathcal{O}(d^3)$ flops which is infeasible for large-scale datasets. Recently, there has been work on solving these problems by leveraging fast linear system solvers [GJK⁺16, AZL17a] while requiring complete knowledge of the matrices A and B .

In the stochastic setting, the difficulty increases because the objective is to maximize a ratio of expectations, in contrast to the standard setting of stochastic optimization [RM51], where the objective is the maximization of an expectation. There has been recent interest in understanding and developing efficient algorithms with provable convergence guarantees for such non-convex problems. [JJK⁺16] and [Sha16] recently analyzed the convergence rate of Oja's algorithm [Oja82], one of the most commonly used algorithm for streaming PCA.

In contrast, for the stochastic generalized eigenvalue problem and CCA problem, the focus has been to translate algorithms from the offline setting to the online one. For example, [GGS⁺17] proposes a streaming algorithm for the stochastic CCA problem which utilizes a streaming SVRG method to solve an online least-squares problem. Despite being streaming in nature, this algorithm requires a non-trivial initialization and, in contrast to the spirit of streaming algorithms, updates its eigenvector estimate only after every few samples. This raises the following challenging question: *Is it possible to obtain an efficient and provably convergent counterpart to*

Oja’s Algorithm for computing the principal generalized eigenvector in the stochastic setting?

In this paper, we propose a simple, globally convergent, *two-line* algorithm, TTS-CCA, for the stochastic principal generalized eigenvector problem and, as a consequence, we obtain a natural extension of Oja’s algorithm for the streaming CCA problem. TTS-CCA is an iterative algorithm which works by updating two coupled sequences at every time step. In contrast with existing methods [JJK⁺16], at each time step the algorithm can be seen as performing a step of Oja’s method, with a noise term which is neither *zero mean* nor *conditionally independent*, but instead is Markovian in nature. The analysis of the algorithm borrows tools from the theory of fast mixing of Markov chains [DDB17] as well as two-time-scale stochastic approximation [BMP90, Bor97, Bor09] to obtain an optimal (up to dimension dependence) fast convergence rate of $\tilde{\mathcal{O}}(1/n)$. Our main contribution can be summarized in the following informal theorem (made formal in Section 5).

Main Result (informal) With probability greater than $4/5$, one can obtain an ϵ -accurate estimate of the generalized eigenvector in the stochastic setting using $\tilde{\mathcal{O}}(1/\epsilon)$ unbiased independent samples of the matrices. The multiplicative pre-factors depend polynomially on the inverse eigengap and the dimension of the problem.

Organization The remainder of this paper is organized as follows. In Section 2, we formalize the problem of streaming CCA and the generalized eigenvalue problem. Section 3 presents the TTS-CCA algorithm and its theoretical analysis. In Section 4, we provide empirical results demonstrating the practical utility of our method, and finally, Section 5 concludes the paper with potential future directions.

Notation We denote by $\lambda_i(M)$ and $\sigma_i(M)$ the i^{th} largest eigenvalue and singular value of a square matrix M . For any positive semi-definite matrix N , we denote inner product in the N -norm by $\langle \cdot, \cdot \rangle_N$ and the corresponding norm by $\|\cdot\|_N$. We let $\kappa_N = \frac{\lambda_{\max}(N)}{\lambda_{\min}(N)}$ denote the condition number of N . We denote the eigenvalues of the matrix $B^{-1}A$ by $\lambda_1 > \lambda_2 \geq \dots \geq \lambda_d$ with $(u_i)_{i=1}^d$ and $(\tilde{u}_i)_{i=1}^d$ denoting the corresponding right and left eigenvectors of $B^{-1}A$ whose existence is guaranteed by Lemma 24 in Appendix G.3. We use Δ_λ to denote the eigengap $\lambda_1 - \lambda_2$.

2 Problem Statement

In this section, we focus on the problem of estimating principal generalized eigenvectors in a stochastic setting. The generalized eigenvector, v_i , corresponding to a system of matrices (A, B) , where $A \in \mathbb{R}^{d \times d}$ is a symmetric matrix and $B \in \mathbb{R}^{d \times d}$ is a symmetric positive definite matrix, satisfies

$$Av_i = \lambda_i Bv_i \tag{4}$$

The principal generalized eigenvector v_1 corresponds to the vector with the largest value¹ of λ_i , or, equivalently, v_1 is the principal eigenvector of the non-symmetric matrix $B^{-1}A$. The vector v_1 also corresponds to the maximizer of the generalized Rayleigh quotient given by

$$v_1 = \arg \max_{v \in \mathbb{R}^d} \frac{v^\top A v}{v^\top B v} \tag{5}$$

¹Note that we consider here the largest *signed* value of λ_i

In the stochastic setting, we only have access to a sequence of matrices $A_1, \dots, A_n \in \mathbb{R}^{d \times d}$ and $B_1, \dots, B_n \in \mathbb{R}^{d \times d}$ assumed to be drawn i.i.d. from an unknown underlying distribution, such that $\mathbb{E}[A_i] = A$ and $\mathbb{E}[B_i] = B$ and the objective is to estimate v_1 given access to $\mathcal{O}(d)$ memory.

In order to quantify the error between a vector and its estimate, we define the following generalization of the sine with respect to the B -norm as,

$$\sin_B^2(v, w) = 1 - \left(\frac{v^\top B w}{\|v\|_B \|w\|_B} \right)^2 \quad (6)$$

3 Related Work

PCA. There is a vast literature dedicated to the development of computationally efficient algorithms for the PCA problem in the offline setting (see [MM15, GHJ⁺16] and references therein). In the stochastic setting, sharp convergence results were obtained recently by [JJK⁺16] and [Sha16] for the principal eigenvector computation problem using Oja’s algorithm and later extended to the streaming k-PCA setting by [AZL17b]. They are able to obtain a $\mathcal{O}(1/n)$ convergence rate when the eigengap of the matrix is positive and a $\mathcal{O}(1/\sqrt{n})$ rate is attained in the gap free setting.

Offline CCA and generalized eigenvector. Computationally efficient optimization algorithms with finite convergence guarantees for CCA and the generalized eigenvector problem based on Empirical Risk Minimization (ERM) on a fixed dataset have recently been proposed in [GJK⁺16, WWGS16, AZL17a]. These approaches work by reducing the CCA and generalized eigenvector problem to that of solving a PCA problem on a modified matrix M (e.g., for CCA, $M = B^{-\frac{1}{2}} A B^{-\frac{1}{2}}$). This reformulation is then solved by using an approximate version of the Power Method that relies on a linear system solver to obtain the approximate power method step. [GJK⁺16, AZL17a] propose an algorithm for the generalized eigenvector computation problem and instantiate their results for the CCA problem. [LF14, MLF15, WWGS16] focus on the CCA problem by optimizing a different objective:

$$\min \frac{1}{2} \widehat{\mathbb{E}} |\phi^\top x_i - \psi^\top y_i|^2 + \lambda_x \|\phi\|_2^2 + \lambda_y \|\psi\|_2^2 \quad \text{s.t.} \quad \|\phi\|_{\widehat{\mathbb{E}}[xx^\top]} = \|\psi\|_{\widehat{\mathbb{E}}[yy^\top]} = 1$$

where $\widehat{\mathbb{E}}$ denotes the empirical expectation. The proposed methods utilize the knowledge of complete data in order to solve the ERM problem, and hence is unclear how to extend them to the stochastic setting.

Stochastic CCA and generalized eigenvector. There has been a dearth of work for solving these problems in the stochastic setting owing to the difficulties mentioned in Section 1. Recently, [GGS⁺17] extend the algorithm of [WWGS16] from the offline to the streaming setting by utilizing a streaming version of the SVRG algorithm for the least squares system solver. Their algorithm, based on the shift and invert method, suffers from two drawbacks: a) contrary to the spirit of streaming algorithms, this method does not update its estimate at each iteration – it requires to use logarithmic samples for solving an online least squares problem, and, b) their algorithm critically relies on obtaining an estimate of λ_1 to a small accuracy for which it requires to burn a few samples in the process. In comparison, TTS-CCA takes a *single* stochastic gradient step for the inner least squares problem and updates its estimate of the eigenvector after each sample. Perhaps the closest to our approach is [AMMS17], who propose an online method by solving

Algorithm 1: TTS-CCA for Streaming $Av = \lambda Bv$

Input: Time steps T , step size α_t (Least Squares), β_t (Oja)

Initialize: $(w_0, v_0) \leftarrow$ sample uniformly from the unit sphere in \mathbb{R}^d , $\bar{v}_0 = v_0$

for $t = 1, \dots, T$ **do**

 Draw sample (A_t, B_t)

$w_t \leftarrow w_{t-1} - \alpha_t(B_t w_{t-1} - A_t v_{t-1})$

$v'_t \leftarrow v_{t-1} + \beta_t w_t$

$v_t \leftarrow \frac{v'_t}{\|v'_t\|_2}$

Output: Estimate of Principal Generalized Eigenvector: v_T

a convex relaxation of the CCA objective with an inexact stochastic mirror descent algorithm. Unfortunately, the computational complexity of their method is $\mathcal{O}(d^2)$ which renders it infeasible for large-scale problems.

4 TTS-CCA

In this section, we describe our proposed approach for the stochastic generalized eigenvector problem (see Section 2). Our algorithm TTS-CCA, described in Algorithm 1, is a natural extension of the popular Oja’s algorithm used for solving the streaming PCA problem. The algorithm proceeds by iteratively updating two coupled sequences (w_t, v_t) at the same time: w_t is updated using one step of stochastic gradient descent with constant step-size to minimize $w^\top Bw - 2w^\top Av_t$ and v_t is updated using a step of Oja’s algorithm. TTS-CCA has its roots in the theory of two-time-scale stochastic approximation, by viewing the sequence w_t as a fast mixing Markov chain and v_t as a slowly evolving one. In the sequel, we describe the evolution of the Markov chains $(w_t)_{t \geq 0}, (v_t)_{t \geq 0}$, in the process outlining the intuition underlying TTS-CCA and understanding the key challenges which arise in the convergence analysis.

Oja’s algorithm. TTS-CCA is closely related to the Oja’s algorithm [Oja82] for the streaming PCA problem. Consider a special case of the problem, when each $B_t = I$. In the offline setting, this reduces the generalized eigenvector problem to that of computing the principal eigenvector of A . With the setting of step-size $\alpha_t = 1$, TTS-CCA recovers the Oja’s algorithm given by

$$v_t = \frac{v_{t-1} + \beta_t A_t v_{t-1}}{\|v_{t-1} + \beta_t A_t v_{t-1}\|}$$

This algorithm is exactly a projected stochastic gradient ascent on the Rayleigh quotient $v^\top Av$ (with a step size β_t). Alternatively, it can be interpreted as a randomized power method on the matrix $(I + \beta_t A)$ [HP14].

Two-time-scale approximation. The theory of two-time-scale approximation forms the underlying basis for TTS-CCA. It considers coupled iterative systems where one component changes much faster than the other [Bor97, Bor09]. More precisely, its objective is to understand classical systems of the type:

$$x_t = x_{t-1} + \alpha_t [h(x_{t-1}, y_{t-1}) + \xi_t^1] \tag{7}$$

$$y_t = y_{t-1} + \beta_t [g(x_{t-1}, y_{t-1}) + \xi_t^2] \quad (8)$$

where g and h are the update functions and (ξ_t^1, ξ_t^2) correspond to the noise vectors at step t and typically assumed to be martingale difference sequences.

In the above model, whenever the two step sizes α_t and β_t satisfy $\beta_t/\alpha_t \rightarrow 0$, the sequence y_t moves on a slower timescale than x_t . For any fixed value of y the dynamical system given by x_t ,

$$x_t = x_{t-1} + \alpha_t [h(x_{t-1}, y) + \xi_t^1] \quad (9)$$

converges to a solution $x^*(y)$. In the coupled system, since the state variables x_t move at a much faster time scale, they can be seen as being close to $x^*(y_t)$, and thus, we can alternatively consider:

$$y_t = y_{t-1} + \beta_t [g(x^*(y_{t-1}), y_{t-1}) + \xi_t^2] \quad (10)$$

If the process given by y_t above were to converge to y^* , under certain conditions, we can argue that the coupled process (x_t, y_t) converges to $(x^*(y^*), y^*)$. Intuitively, because x_t and y_t are evolving at different time-scales, x_t views the process y_t as quasi-constant while y_t views x_t as a process rapidly converging to $x^*(y_t)$.

TTS-CCA can be seen as a particular instance of the coupled iterative system given by Equations (7) and (8) where the sequence v_t evolves with a step-size $\beta_t \approx \frac{1}{t}$, much slower than the sequence w_t , which has a step-size of $\alpha_t \approx \frac{1}{\log(t)}$. Proceeding as above, the sequence v_t views w_t as having converged to $B^{-1}Av_t + \xi_t$, where ξ_t is a noise term, and the update step for v_t in TTS-CCA can be viewed as a step of Oja's algorithm, albeit with Markovian noise.

While previous works on the stochastic CCA problem required to use logarithmic independent samples to solve the inner least-squares problem in order to perform an approximate power method (or Oja) step, the theory of two-time-scale stochastic approximation suggests that it is possible to obtain a similar effect by evolving the sequences w_t and v_t at two different time scales.

Understanding the Markov Process $\{w_t\}$. In order to understand the process described by the sequence w_t , we consider the homogeneous Markov chain (w_t^v) defined by

$$w_t^v = w_{t-1}^v - \alpha(B_t w_{t-1}^v - A_t v) \quad (11)$$

for a constant vector v and we denote its t -step kernel by π_v^t [MT09]. This Markov process is an iterative linear model and has been extensively studied by [Ste99, DF99, BM13]. It is known that for any step-size $\alpha \leq 2/R^2$, the Markov chain $(w_t^v)_{t \geq 0}$ admits a unique stationary distribution, denoted by ν_v . In addition,

$$W_2^2(\pi_v^t(w_0, \cdot), \nu_v) \leq (1 - 2\mu\alpha(1 - \alpha R_B^2/2))^t \int_{\mathbb{R}^d} \|w_0 - w\|_2^2 d\nu_v(w) \quad (12)$$

where $W_2^2(\lambda, \nu)$ denotes the Wasserstein distance of order 2 between probability measures λ and ν (see, e.g., [Vil08] for more properties of W_2). Equation (12) implies that the iterative linear process described by (11) mixes exponentially fast to the stationary distribution. This forms a crucial ingredient in our convergence analysis where we use the fast mixing to obtain a bound on the expected norm of the Markovian noise (see Lemma 1).

Moreover, one can compute the mean \bar{w}^v of the process w_t under the stationary distribution by taking expectation under ν_v on both sides in equation (11). Doing so, we obtain, $\bar{w}^v = B^{-1}Av$. Thus, in our setting, since the v_t process evolves slowly, we can expect that $w_t \approx B^{-1}Av_t$, allowing TTS-CCA to mimic Oja's algorithm.

5 Main Theorem

In this section, we present our main convergence guarantee for TTS-CCA when applied to the streaming generalized eigenvector problem. We begin by listing the key assumptions required by our analysis:

- (A1) The matrices $(A_i)_{i \geq 0}$ satisfy $\mathbb{E}[A_i] = A$ for a symmetric matrix $A \in \mathbb{R}^{d \times d}$.
- (A2) The matrices $(B_i)_{i \geq 0}$ are such that each $B_i \succcurlyeq 0$ is symmetric and satisfies $\mathbb{E}[B_i] = B$ for a symmetric matrix $B \in \mathbb{R}^{d \times d}$ with $B \succcurlyeq \mu I$ for $\mu > 0$.
- (A3) There exists $R \geq 0$ such that $\max\{\|A_i\|, \|B_i\|\} \leq R$ almost surely.

Under the assumptions stated above, we obtain the following convergence theorem for TTS-CCA with respect to the \sin_B^2 distance, as described in Section 2.

Theorem 1 (Main Result). *Fix any $\delta > 0$ and $\epsilon_1 > 0$. Suppose that the step sizes are set to $\alpha_t = \frac{c}{\log(d^2\beta+t)}$ and $\beta_t = \frac{\gamma}{\Delta_\lambda(d^2\beta+t)}$ for $\gamma > 1/2$, $c > 1$ and*

$$\beta = \max \left(\frac{20\gamma^2\lambda_1^2}{\Delta_\lambda^2 d^2 \log \left(\frac{1+\delta/100}{1+\epsilon_1} \right)}, \frac{200 \left(\frac{R}{\mu} + \frac{R^3}{\mu^2} + \frac{R^5}{\mu^3} \right) \log \left(1 + \frac{R^2}{\mu} + \frac{R^4}{\mu^2} \right)}{\delta \Delta_\lambda^2} \right)$$

Suppose that the number of samples n satisfy

$$\frac{d^2\beta + n}{\log^{\frac{1}{\min(1, 2\gamma\lambda_1/\Delta_\lambda)}}(d^2\beta + n)} \geq \left(\frac{cd}{\delta_1 \min(1, \lambda_1)} \right)^{\frac{1}{\min(1, 2\gamma\lambda_1/\Delta_\lambda)}} (d^3\beta + 1) \exp \left(\frac{c\lambda_1^2}{d^2} \right)$$

Then, the output v_n of Algorithm 1 satisfies,

$$\sin_B^2(u_1, v_n) \leq \frac{(2 + \epsilon_1)cd \left\| \sum_{i=1}^d \tilde{u}_i \tilde{u}_i^\top \right\|_2 \log \left(\frac{1}{\delta} \right)}{\delta^2 \|\tilde{u}_1\|_2^2} \left(\frac{c\gamma^2 \log^3(d^2\beta + n)}{\Delta_\lambda^2(d^2\beta + n + 1)} + \frac{cd}{\Delta_\lambda} \left(\frac{d^2\beta + \log^3(d^2\beta)}{d^2\beta + n + 1} \right)^{2\gamma} \right)$$

with probability at least $1 - \delta$ with c depending polynomially on parameters of the problem $\lambda_1, \kappa_B, R, \mu$. The parameter δ_1 is set as $\delta_1 = \frac{\epsilon_1}{2(2+\epsilon_1)}$.

The above result shows that with probability at least $1 - \delta$, TTS-CCA converges in the B -norm to the right eigenvector, u_1 , corresponding to the maximum eigenvalue of the matrix $B^{-1}A$. Further, TTS-CCA exhibits an $\tilde{\mathcal{O}}(1/n)$ rate of convergence, which is known to be optimal for stochastic approximation algorithms even with convex objectives [NY83].

Comparison with Streaming PCA. In the setting where $B = I$, and $A \succeq 0$ is a covariance matrix, the principal generalized eigenvector problem reduces to performing PCA on the A . When compared with the results obtained for streaming PCA by [JJK⁺16], our corresponding results differ by a factor of dimension d and problem dependent parameters $\lambda_1, \Delta_\lambda$. We believe that such a dependence is not inherent to TTS-CCA but a consequence of our analysis. We leave this task of obtaining a dimension free bound for TTS-CCA as future work.

Gap-independent step size : While the step size for the sequence v_n in TTS-CCA depends on eigen-gap, which is a priori unknown, one can leverage recent results as in [TFBJ18] to get around this issue by using a streaming average step size.

6 Proof Sketch

In this section, we detail out the two key ideas underlying the analysis of TTS-CCA to obtain the convergence rate mentioned in Theorem 1: a) controlling the non i.i.d. Markovian noise term which is introduced because of the coupled Markov chains in TTS-CCA and b) proving that a noisy power method with such Markovian noise converges to the correct solution.

Controlling Markovian perturbations. In order to better understand the sequence v_t , we rewrite the update as,

$$v'_t = v_{t-1} + \beta_t w_t = v_{t-1} + \beta_t (B^{-1} A v_{t-1} + \xi_t) \quad (13)$$

where $\xi_t = w_t - B^{-1} A v_{t-1}$ is the prediction error which is a Markovian noise. Note that the noise term is neither *mean zero* nor a *martingale difference* sequence. Instead, the noise term ξ_t is dependent on all previous iterates, which makes the analysis of the process more involved. This framework with Markovian noise has been extensively studied by [BMP90, AMP05].

From the update in Equation (13), we observe that TTS-CCA is performing an Oja update but with a controlled Markovian noise. However, we would like to highlight that classical techniques in the study of stochastic approximation with Markovian noise (as the *Poisson Equation* [BMP90, MT09]) were not enough to provide adequate control on the noise to show convergence.

In order to overcome this difficulty, we leverage the fast mixing of the chain w_t^v for understanding the Markovian noise. While it holds that $\mathbb{E}[\|\xi_t\|_2] = \mathcal{O}(1)$ (see Appendix C), a key part of our analysis is the following lemma, the proof of which can be found in Appendix B).

Lemma 1. . For any choice of $k > 4 \frac{\lambda_1(B)}{\mu\alpha} \log(\frac{1}{\beta_{t+k}})$, and assuming that $\|w_s\| \leq W_s$ for $t \leq s \leq t+k$ we have that

$$\|\mathbb{E}[\xi_{t+k} | \mathcal{F}_t]\|_2 = \mathcal{O}(\beta_t k^2 \alpha_t W_{t+k})$$

Lemma 1 uses the fast mixing of w_t to show that $\|\mathbb{E}[\xi_t] | \mathcal{F}_{t-r}\|_2 = \tilde{\mathcal{O}}(\beta_t)$ where $r = \mathcal{O}(\log t)$, i.e., the magnitude of the expected noise is small conditioned on $\log(t)$ steps in the past.

Analysis of Oja's algorithm. The usual proofs of convergence for stochastic approximation define a Lyapunov function and show that it decreases sufficiently at each iteration. Oftentimes control on the per step rate of decrease can then be translated into a global convergence result. Unfortunately in the context of PCA, due to the non-convexity of the Raleigh quotient, the quality of the estimate v_t cannot be related to the previous v_{t-1} . Indeed v_t may become orthogonal to the leading eigenvector. Instead [JJK⁺16] circumvent this issue by leveraging the randomness of the initialization and adopt an operator view of the problem. We take inspiration from this approach in our analysis of TTS-CCA. Let $G_i = w_i v_{i-1}^\top$ and $H_t = \prod_{i=1}^t (I + \beta_i G_i)$, TTS-CCA's update can be equivalently written as

$$v_t = \frac{H_t v_0}{\|H_t v_0\|_2^2}$$

pushing, for the analysis only, the normalization step at the end. This point of view enables us to analyze the improvement of H_t over H_{t-1} since allows one to interpret Oja's update as one step of power method on H_t starting on a random vector v_0 . We present here an easy adaptation of [JJK⁺16, Lemma 3.1] that takes into account the special geometry of the generalized eigenvector problem and the asymmetry of $B^{-1}A$. The proof can be found in Appendix A.

Lemma 2. Let $H \in \mathbb{R}^{d \times d}$, $(u_i)_{i=1}^d$ and $(\tilde{u}_i)_{i=1}^d$ be the corresponding right and left eigenvectors of $B^{-1}A$ and $w \in \mathbb{R}^d$ chosen uniformly on the sphere, then with probability $1 - \delta$ (over the randomness in the initial iterate)

$$\sin_B^2(u_i, Hw) \leq \frac{C \log(1/\delta)}{\delta} \frac{\text{Tr}(HH^\top \sum_{j \neq i} \tilde{u}_j \tilde{u}_j^\top)}{\tilde{u}_i^\top HH^\top \tilde{u}_i} \quad (14)$$

for some universal constant $C > 0$.

This lemma has the virtue of highly simplifying the challenging proof of convergence of Oja's algorithm. Indeed we only have to prove that H_t will be close to $\prod_{i=1}^t (I + \beta_i B^{-1}A)$ for t large enough which can be interpreted as an analogue of the law of large numbers for the multiplication of matrices. This will ensure that $\text{Tr}(H_t H_t^\top \sum_{j \neq i} \tilde{u}_j \tilde{u}_j^\top)$ is relatively small compared to $\tilde{u}_i^\top H_t H_t^\top \tilde{u}_i$ and be enough with Lemma 2 to prove Theorem 1. The proof follows the line of [JJK⁺16] with two additional tedious difficulties: the Markovian noise is neither unbiased nor independent of the previous iterates, and the matrix $B^{-1}A$ is no longer symmetric, which is precisely why we consider the left eigenvector \tilde{u}_i in the right-hand side of Eq. (14). We highlight two key steps:

- First we show that $\mathbb{E} \text{Tr}(H_t H_t^\top \sum_{j \neq i} \tilde{u}_j \tilde{u}_j^\top)$ grows as $\mathcal{O}(\exp(2\lambda_2 \sum_{i=1}^t \beta_i))$, which implies by Markov's inequality the same bound on $\text{Tr}(H_t H_t^\top \sum_{j \neq i} \tilde{u}_j \tilde{u}_j^\top)$ with constant probability. See Lemmas 16 for more details.
- Second we show that $\text{Var} \tilde{u}_i^\top H_t H_t^\top \tilde{u}_i$ grows as $\mathcal{O}(\exp(4\lambda_1 \sum_{i=1}^t \beta_i))$ and $\mathbb{E} \tilde{u}_i^\top HH^\top \tilde{u}_i$ grows as $\mathcal{O}(\exp(2\lambda_1 \sum_{i=1}^t \beta_i))$ which implies by Chebyshev's inequality the same bound for $\tilde{u}_i^\top HH^\top \tilde{u}_i$ with constant probability. See Lemmas 17 and 19 for more details.

7 Application to Canonical Correlation Analysis

Consider two random vectors $X \in \mathbb{R}^d$ and $Y \in \mathbb{R}^d$ with joint distribution P_{XY} . The objective of canonical correlation analysis in the population setting is to find the canonical correlation vectors $\phi, \psi \in \mathbb{R}^{d,d}$ which maximize the correlation

$$\max_{\phi, \psi} \frac{\mathbb{E}[(\phi^\top X)(\psi^\top Y)]}{\sqrt{\mathbb{E}[(\phi^\top X)^2] \mathbb{E}[(\psi^\top Y)^2]}}$$

This problem is equivalent to maximizing $\phi^\top \mathbb{E}[XY^\top] \psi$ under the constraint $\mathbb{E}[(\phi^\top X)^2] = \mathbb{E}[(\psi^\top Y)^2] = 1$ and admits a closed form solution: if we define $T = \mathbb{E}[XX^\top]^{-1/2} \mathbb{E}[XY^\top] \mathbb{E}[YY^\top]^{-1/2}$, then the solution is $(\phi_*, \psi_*) = (\mathbb{E}[XX^\top]^{-1/2} a_1 \mathbb{E}[YY^\top]^{-1/2} b_1)$ where a_1, b_1 are the left and right principal singular vectors of T . By the KKT conditions, there exist $\nu_1, \nu_2 \in \mathbb{R}$ such that this solution satisfies the stationarity equation

$$\mathbb{E}[XY^\top] \psi = \nu_1 \mathbb{E}[XX^\top] \phi \quad \text{and} \quad \mathbb{E}[YX^\top] \phi = \nu_2 \mathbb{E}[YY^\top] \psi$$

Using the constraint conditions we conclude that $\nu_1 = \nu_2$. This condition can be written (for $\lambda = \nu_1$) in the matrix form of Eq. (3). As a consequence, finding the largest generalized eigenvector for the matrices (A, B) will recover the canonical correlation vector (ϕ, ψ) . Solving the associated generalized streaming eigenvector problem, we obtain the following result for estimating the canonical correlation vector whose proof easily follows from Theorem 1 (setting $\gamma = 6$).




Figure 1. Synthetic Generalized Eigenvalue problem. Left: Comparison with two-steps methods. Middle: Robustness to step size α_t . Right: Robustness to step size β_t (Streaming averaged TTS-CCA is dashed).

Theorem 2. Assume that $\max\{\|X\|, \|Y\|\} \leq R$ a.s., $\min\{\lambda_{\min}(\mathbb{E}[XX^\top]), \lambda_{\min}(\mathbb{E}[YY^\top])\} = \mu > 0$ and $\sigma_1(T) - \sigma_2(T) = \Delta > 0$. Fix any $\delta > 0$, let $\epsilon_1 \geq 0$, and suppose the step sizes are set to $\alpha_t = \frac{1}{2R^2 \log(d^2\beta+t)}$ and $\beta_t = \frac{6}{\Delta(d^2\beta+t)}$ and

$$\beta = \max \left(\frac{720\sigma_1^2}{\Delta^2 d^2 \log \left(\frac{1+\delta/100}{1+\epsilon_1} \right)}, \frac{200 \left(\frac{R}{\mu} + \frac{R^3}{\mu^2} + \frac{R^5}{\mu^3} \right) \frac{1}{\delta} \log \left(1 + \frac{R^2}{\mu} + \frac{R^4}{\mu^2} \right)}{\Delta^2} \right)$$

Suppose that the number of samples n satisfy

$$\frac{d^2\beta + n}{\log^{\frac{1}{\min(1, 12\lambda_1/\Delta\lambda)}}(d^2\beta + n)} \geq \left(\frac{cd}{\delta_1 \min(1, \lambda_1)} \right)^{\frac{1}{\min(1, 12\lambda_1/\Delta\lambda)}} (d^3\beta + 1) \exp \left(\frac{c\lambda_1^2}{d^2} \right)$$

Then the output (ϕ_t, ψ_t) of Algorithm 1 applied to (A, B) defined above satisfies,

$$\sin_B^2((\phi_*, \psi_*), (\phi_t, \psi_t)) \leq \frac{(2 + \epsilon_1)cd^2 \log \left(\frac{1}{\delta} \right)}{\delta^2 \|\tilde{u}_1\|_2^2} \frac{\log^3(d^2\beta + n)}{\Delta^2(d^2\beta + n + 1)}$$

with probability at least $1 - \delta$ with c depending on parameters of the problem and independent of d and Δ where $\delta_1 = \frac{\epsilon_1}{2(2+\epsilon_1)}$.

We can make the following observations:

- The convergence guarantee are comparable with the sample complexity obtained by the ERM ($t = \tilde{\mathcal{O}}(d/(\epsilon\Delta^2))$ for sub-Gaussian variables and $t = \tilde{\mathcal{O}}(1/(\epsilon\Delta^2\mu^2))$ for bounded variables)[GGS⁺17] and matches the lower bound $t = \mathcal{O}(d/(\epsilon\Delta^2))$ known for sparse CCA [GMZ17].
- The sample complexity in [GGS⁺17] is better in term of the dependence on d . They obtain the same rates as the ERM. The comparison with [AMMS17] is meaningless since they are in the gap free setting and their computational complexity is $\mathcal{O}(d^2)$.

8 Simulations

Here we illustrate the practical utility of TTS-CCA on a synthetic, streaming generalized eigenvector problem. We take $d = 20$ and $T = 10^6$. The streams $(A_t, B_t) \in (\mathbb{R}^{d \times d})^2$ are normally-distributed with covariance matrix A and B with random eigenvectors and eigenvalues decaying as $1/i$, for $i = 1, \dots, d$. Here R^2 denotes the radius of the streams with $R^2 = \max\{\text{Tr } A, \text{Tr } B\}$. All results are averaged over ten repetitions.

Comparison with two-steps methods. In the left plot of Figure 1 we compare the behavior of TTS-CCA to different two-steps algorithms. Since the method by [AMMS17] is of complexity $\mathcal{O}(d^2)$, we compare TTS-CCA to a method which alternates between one step of Oja’s algorithm and τ steps of averaged stochastic gradient descent with constant step size $1/2R^2$. TTS-CCA is converging at rate $\mathcal{O}(1/t)$ whereas the other methods are very slow. For $\tau = 10$, the solution of the inner loop is too inaccurate and the steps of Oja are inefficient. For $\tau = 10000$, the output of the sgd steps is very accurate but there are too few Oja iterations to make any progress. $\tau = 1000$ seems an optimal parameter choice but this method is slower than TTS-CCA by an order of magnitude.

Robustness to incorrect step-size α . In the middle plot of Figure 1 we compare the behavior of TTS-CCA for step size $\alpha \in \{\alpha_*, \alpha_*/8, \alpha_*/16\}$ where $\alpha_* = 1/R^2$. We observe that TTS-CCA converges at a rate $\mathcal{O}(1/t)$ independently of the choice of α .

Robustness to incorrect step-size β_t . In the right plot of Figure 1 we compare the behavior of TTS-CCA for step size $\beta_t \in \{\beta_*/t, \beta_*/16t, \beta_*/\sqrt{t}, \beta_*/16\sqrt{t}\}$ where β_* corresponds to the minimal error after one pass over the data. We observe that TTS-CCA is not robust to the choice of the constant for step size $\beta_t \propto 1/t$. If the constant is too small, the rate of convergence is arbitrary slow. We observe that considering the streaming average of [TFBJ18] on TTS-CCA with a step size $\beta_t \propto 1/\sqrt{t}$ enables to recover the fast $\mathcal{O}(1/t)$ convergence while being robust to constant misspecification.

9 Conclusion

In this paper, we tackle the problems of principal Generalized Eigenvector computation and Canonical Correlation Analysis (CCA) in the stochastic setting. We introduce TTS-CCA, a simple and efficient two-time-scale stochastic approximation algorithm that extends Oja’s method to the streaming scenario. Our analysis demonstrates the global convergence of TTS-CCA by leveraging techniques from fast-mixing Markov chains and two-time-scale stochastic approximation. We show that TTS-CCA achieves an optimal convergence rate of $\mathcal{O}(1/n)$, making it highly suitable for large-scale, high-dimensional data.

Furthermore, our development of tools for understanding stochastic processes with Markovian noise provides new insights into algorithmic behavior under such conditions, which may have broader applicability beyond CCA and GEP. Future work may explore the extension of TTS-CCA to more complex eigenvalue and multi-view learning problems, as well as further empirical evaluation on real-world streaming datasets. Theoretical advances such as dimension-free convergence guarantees and gap-independent step sizes are promising areas for future exploration.

References

- [AMMS17] R. Arora, T. V. Marinov, P. Mianjy, and N. Srebro. Stochastic approximation for canonical correlation analysis. In *Advances in Neural Information Processing Systems*. 2017.
- [AMP05] C. Andrieu, E. Moulines, and P. Priouret. Stability of stochastic approximation under verifiable conditions. *SIAM Journal on Control and Optimization*, 44(1):283–312, 2005.
- [AZL17a] Z. Allen-Zhu and Y. Li. Doubly accelerated methods for faster CCA and generalized eigendecomposition. In *International Conference on Machine Learning*, 2017.

- [AZL17b] Zeyuan Allen-Zhu and Yuanzhi Li. First efficient convergence for streaming k-PCA: a global, gap-free, and near-optimal rate. In *Foundations of Computer Science (FOCS), 2017 IEEE 58th Annual Symposium on*, pages 487–492. IEEE, 2017.
- [BM13] F. Bach and E. Moulines. Non-strongly-convex smooth stochastic approximation with convergence rate $O(1/n)$. In *Advances in Neural Information Processing Systems*, 2013.
- [BMP90] A. Benveniste, M. Métivier, and P. Priouret. *Adaptive Algorithms and Stochastic Approximations*. Springer Publishing Company, Incorporated, 1990.
- [Bor97] Vivek S Borkar. Stochastic approximation with two time scales. *Systems & Control Letters*, 29(5):291–294, 1997.
- [Bor09] Vivek S Borkar. *Stochastic Approximation: A Dynamical Systems Viewpoint*, volume 48. Springer, 2009.
- [CKLS09] K. Chaudhuri, S. M Kakade, K. Livescu, and K. Sridharan. Multi-view clustering via canonical correlation analysis. In *International Conference on Machine Learning*, pages 129–136, 2009.
- [DDB17] A. Dieuleveut, A. Durmus, and F. Bach. Bridging the gap between constant step size stochastic gradient descent and Markov chains. *arXiv preprint arXiv:1707.06386*, 2017.
- [DF99] P. Diaconis and D. Freedman. Iterated random functions. *SIAM Review*, 41(1):45–76, 1999.
- [GGS⁺17] Chao Gao, Dan Garber, Nathan Srebro, Jialei Wang, and Weiran Wang. Stochastic canonical correlation analysis. *arXiv preprint arXiv:1702.06533*, 2017.
- [GHJ⁺16] D. Garber, E. Hazan, J. Jin, S. Kakade, C. Musco, P. Netrapalli, and A. Sidford. Faster eigenvector computation via shift-and-invert preconditioning. In *International Conference on Machine Learning*, 2016.
- [GJK⁺16] R. Ge, C. Jin, S. Kakade, P. Netrapalli, and A. Sidford. Efficient algorithms for large-scale generalized eigenvector computation and canonical correlation analysis. In *International Conference on Machine Learning*, 2016.
- [GMZ17] C. Gao, Z. Ma, and H. Zhou. Sparse CCA: Adaptive estimation and computational barriers. *The Annals of Statistics*, 45(5):2074–2101, 2017.
- [Hot36] H. Hotelling. Relations between two sets of variates. *Biometrika*, 28(3/4):321–377, 1936.
- [HP14] M. Hardt and E. Price. The noisy power method: A meta algorithm with applications. In *Advances in Neural Information Processing Systems*, pages 2861–2869, 2014.
- [JJJ⁺16] P. Jain, C. Jin, S. M. Kakade, P. Netrapalli, and A. Sidford. Streaming PCA: Matching matrix Bernstein and near-optimal finite sample guarantees for Oja’s algorithm. In *Conference on Learning Theory*, 2016.
- [KF07] S. M Kakade and D. P Foster. Multi-view regression via canonical correlation analysis. In *International Conference on Computational Learning Theory*, pages 82–96. Springer, 2007.
- [KM13] Nikos Karampatziakis and Paul Mineiro. Discriminative features via generalized eigenvectors. *arXiv preprint arXiv:1310.1934*, 2013.
- [LF14] Y. Lu and D. P. Foster. Large scale canonical correlation analysis with iterative least squares. In *Advances in Neural Information Processing Systems*, 2014.
- [MLF15] Z. Ma, Y. Lu, and D. Foster. Finding linear structure in large datasets with scalable canonical correlation analysis. In *International Conference on Machine Learning*, 2015.
- [MM15] C. Musco and C. Musco. Randomized block Krylov methods for stronger and faster approximate singular value decomposition. In *Advances in Neural Information Processing Systems*, 2015.

- [MT09] S. P. Meyn and R. L. Tweedie. *Markov Chains and Stochastic Stability*. Cambridge University Press, 2009.
- [NY83] A. S. Nemirovsky and D. B. Yudin. *Problem Complexity and Method Efficiency in Optimization*. Wiley-Interscience Series in Discrete Mathematics. John Wiley & Sons, 1983.
- [Oja82] Erkki Oja. Simplified neuron model as a principal component analyzer. *Journal of Mathematical Biology*, 15(3):267–273, 1982.
- [RM51] Herbert Robbins and Sutton Monro. A stochastic approximation method. *The Annals of Mathematical Statistics*, 22(3):400–407, 1951.
- [Sha16] O. Shamir. Convergence of stochastic gradient descent for PCA. In *International Conference on Machine Learning*, 2016.
- [Ste99] D. Steinsaltz. Locally contractive iterated function systems. *Annals of Probability*, pages 1952–1979, 1999.
- [TFBJ18] N. Tripuraneni, N. Flammarion, F. Bach, and M. I. Jordan. Averaging stochastic gradient descent on Riemannian manifolds. In *Conference on Learning Theory*, 2018.
- [Vil08] C. Villani. *Optimal Transport: Old and New*, volume 338. Springer-Verlag Berlin Heidelberg, 2008.
- [WWGS16] W. Wang, J. Wang, D. Garber, and N. Srebro. Efficient globally convergent stochastic optimization for canonical correlation analysis. In *Advances in Neural Information Processing Systems*, 2016.